

CS290D – Advanced Data Mining

Instructor: Xifeng Yan
Computer Science
University of California at Santa Barbara

Recursive Neural Networks

Lecturer: Semih Yavuz
Computer Science
University of California at Santa Barbara

Source of slides

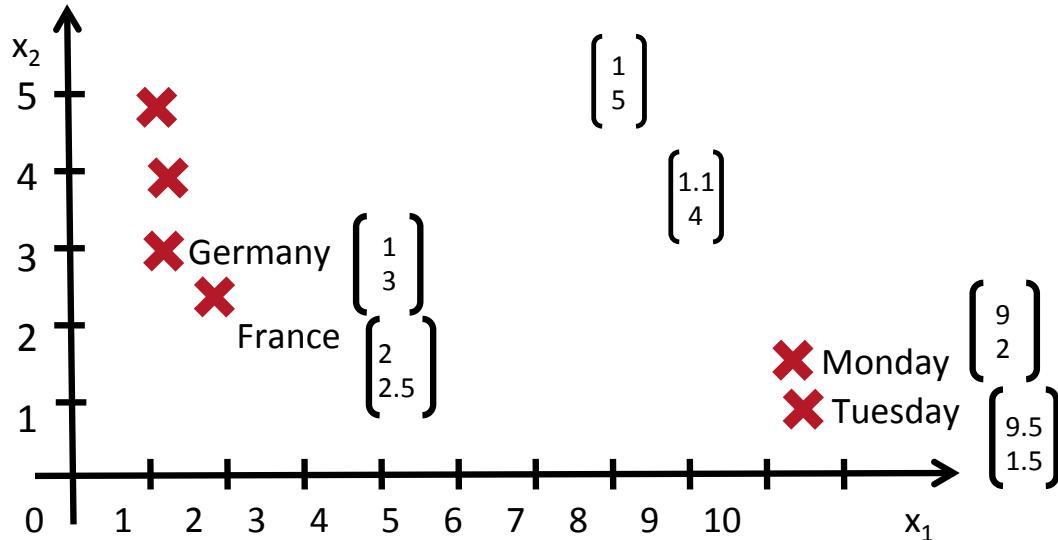
- Deep Learning for Natural Language Processing Course by Richard Socher - 2015 Stanford
- Lecture Notes by Prof. Charles Elkan: [http://cseweb.ucsd.edu/~elkan/250B/
learningmeaning.pdf](http://cseweb.ucsd.edu/~elkan/250B/learningmeaning.pdf)
- Deep Learning for NLP (without Magic) by Richard Socher - NAACL 2013

Outline

- Motivation
- Recursive Neural Networks
 - Compositionality
 - Structure Prediction: Parsing
 - Backpropagation Through Structure
 - Minor twist leads to an improved Parser
- Matrix-Vector RNNs
- Recursive Neural Tensor Networks

Word vectors

- How to represent the meaning of longer phrases?



the country of my birth
the place where I was born



In a way that the phrases of same meaning are close to each other.

Example use: **Paraphrasing, Text Summarization**

- How to map the longer phrases in to the same vector space?

Semantic Vector Space Embeddings

Many works for single word embeddings!



Single Word Vectors

- Distributional Techniques
- Brown Clusters
- Useful as features inside models, e.g. CRFs for NER, etc.
- Cannot capture longer phrases

Ignores Word Order

Documents Vectors

- Bag of words models
- PCA (LSA, LDA)
- Great for IR, document exploration, etc.
- Ignore word order, no detailed understanding

Semantic Vector Space Embeddings

Vectors representing
Phrases and Sentences
that do not ignore word order
and capture semantics for NLP tasks



Single Word Vectors

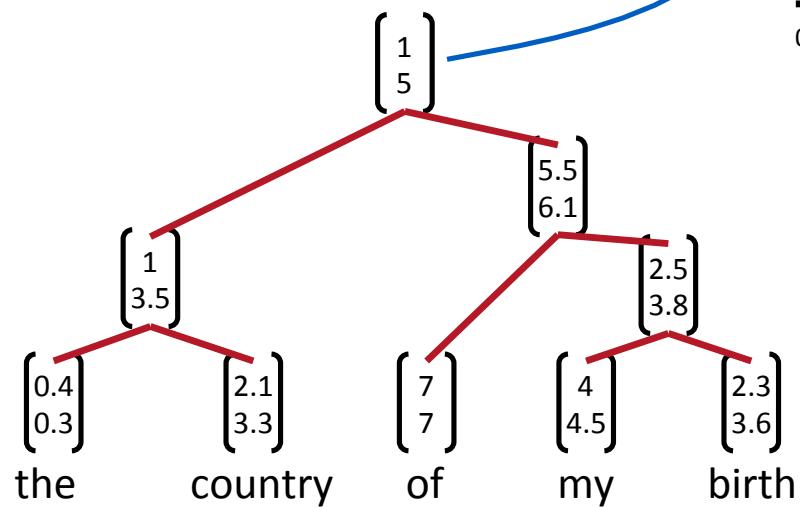
- Distributional Techniques
- Brown Clusters
- Useful as features inside models, e.g. CRFs for NER, etc.
- Cannot capture longer phrases

Documents Vectors

- Bag of words models
- PCA (LSA, LDA)
- Great for IR, document exploration, etc.
- Ignore word order, no detailed understanding

Mapping phrases into a vector space

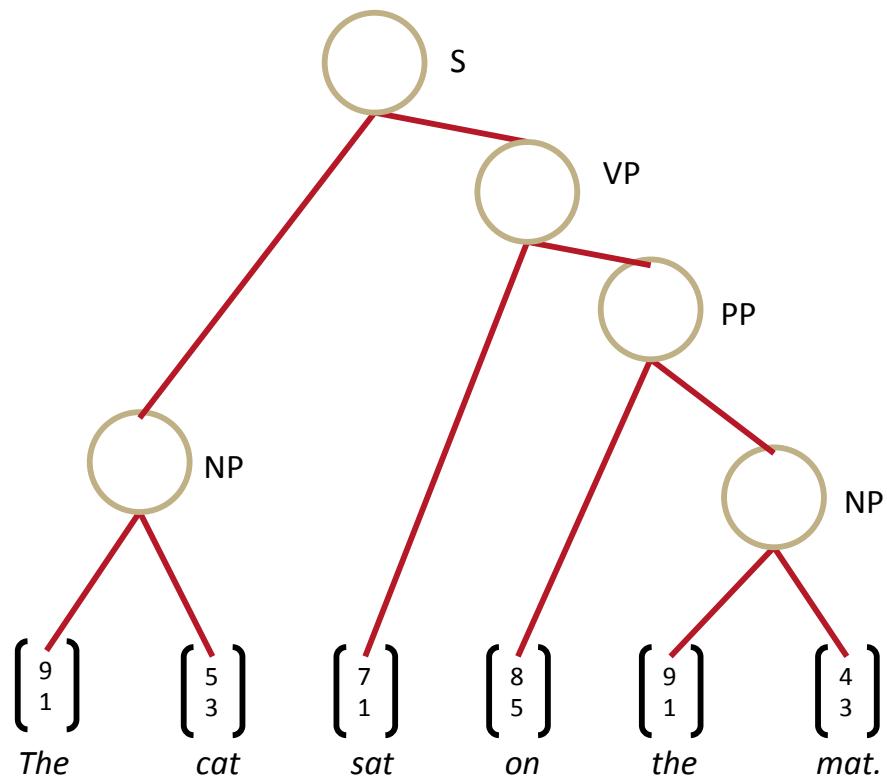
- How to achieve **compositionality**
- Define the meaning vector of phrases recursively by
 - the meaning of its sub phrases (words as atoms)
 - learning the rules that combine them



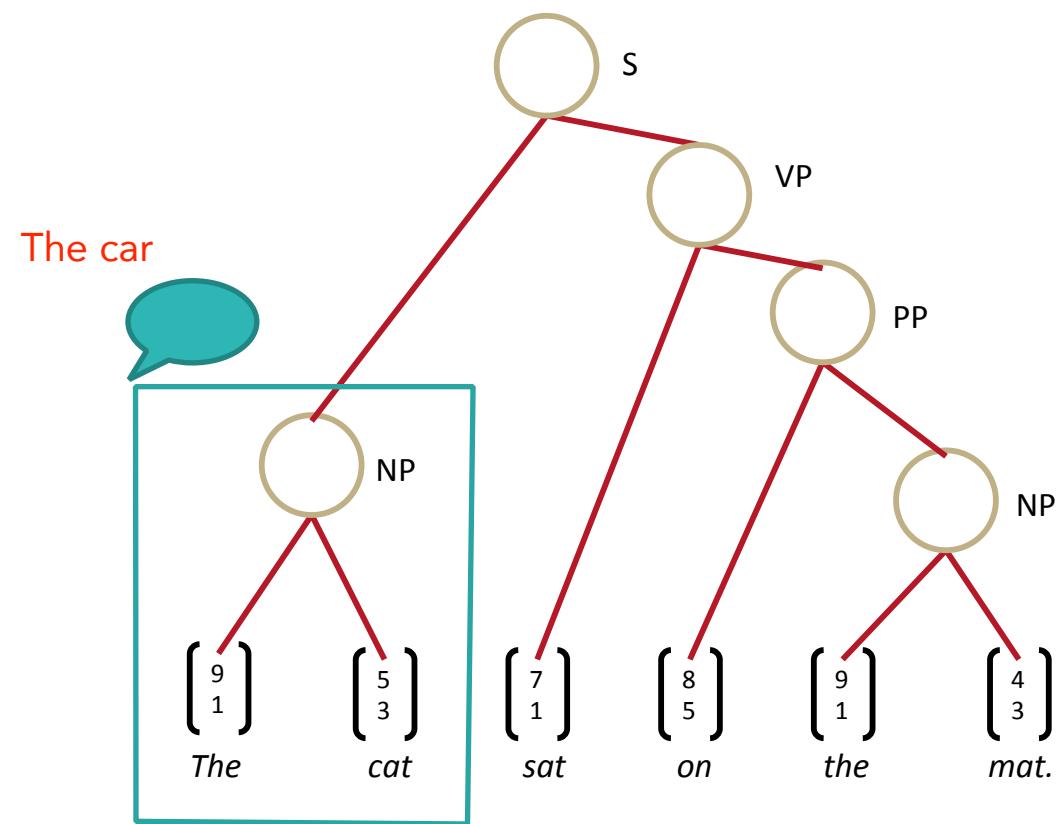
Outline

- Motivation
- Recursive Neural Networks
 - Compositionality
 - Structure Prediction: Parsing
 - Backpropagation Through Structure
 - Minor twist leads to an improved Parser
- Matrix-Vector RNNs
- Recursive Neural Tensor Networks

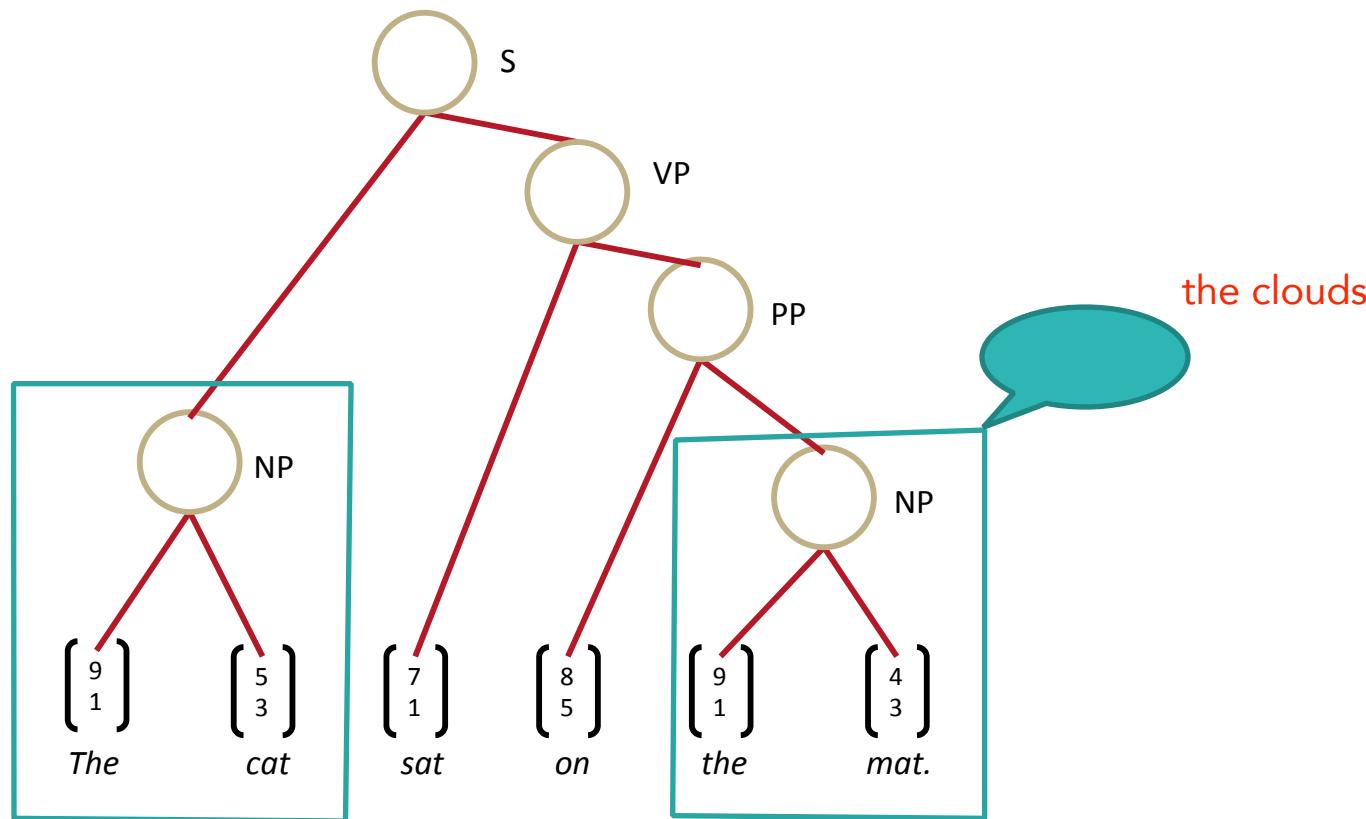
Parsing: What is the objective?



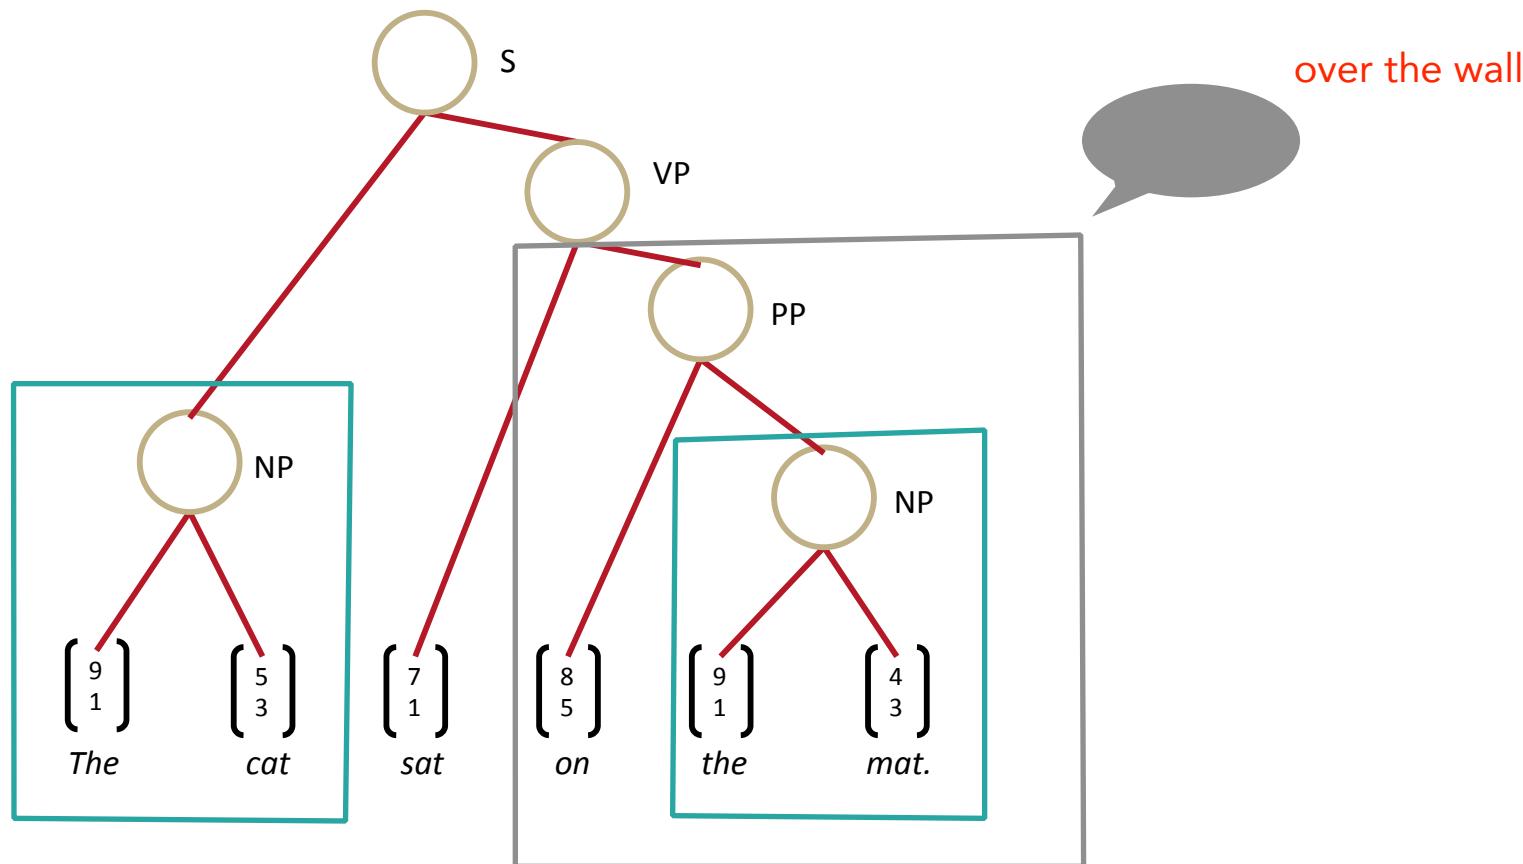
Parsing: What is the objective?



Parsing: What is the objective?

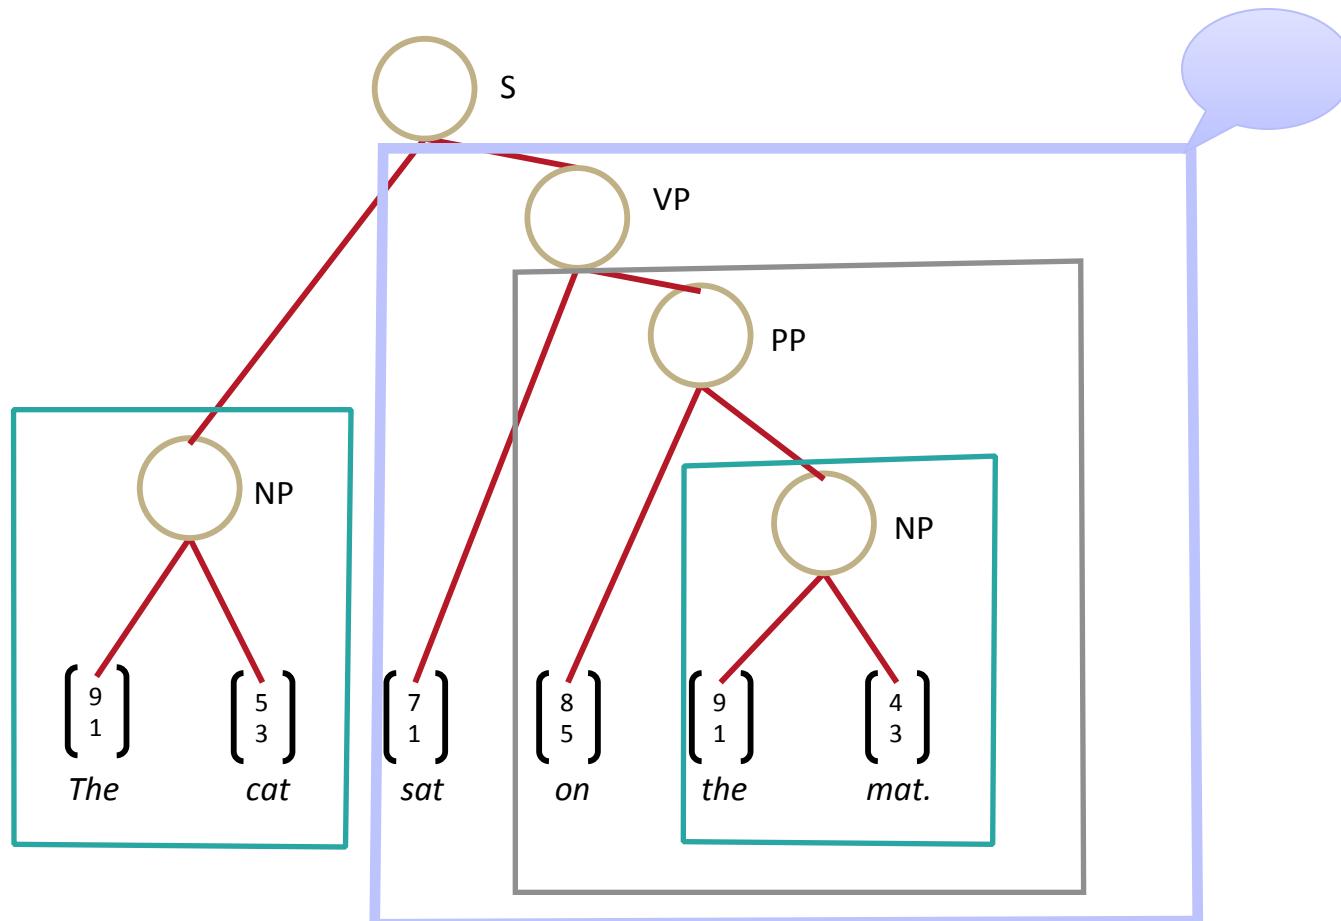


Parsing: What is the objective?

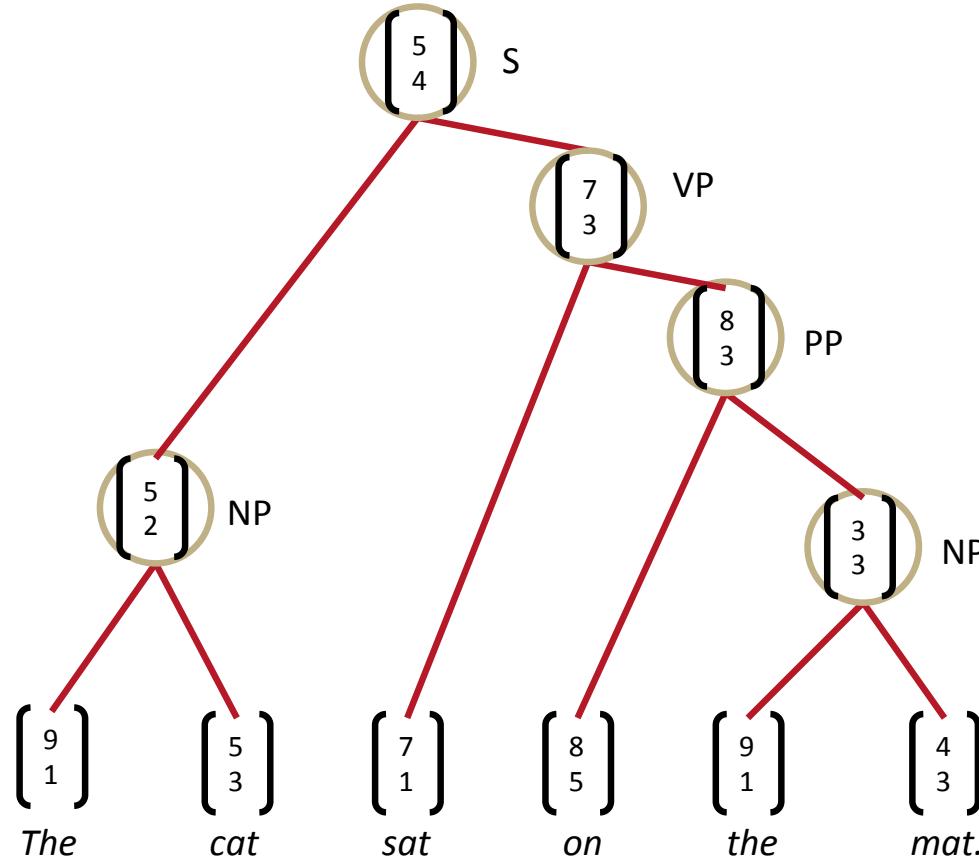


Parsing: What is the objective?

walked into the room



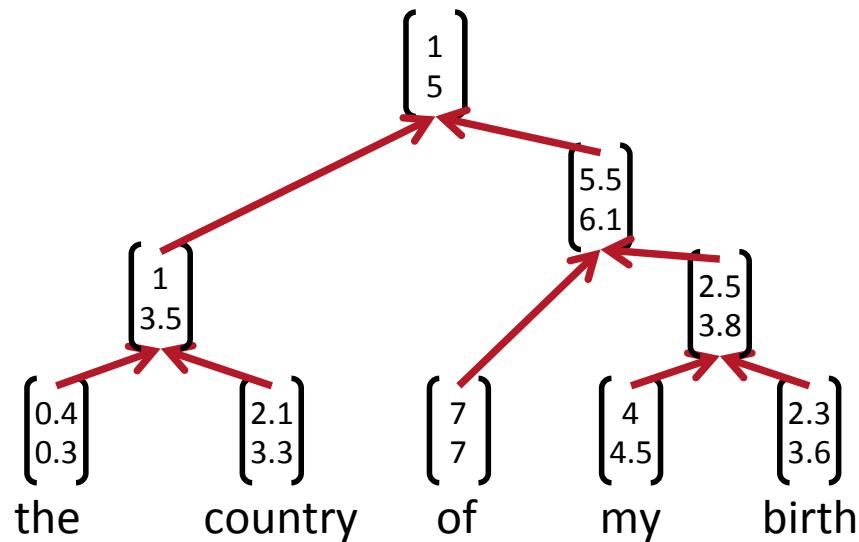
Learn jointly Structure and Representation



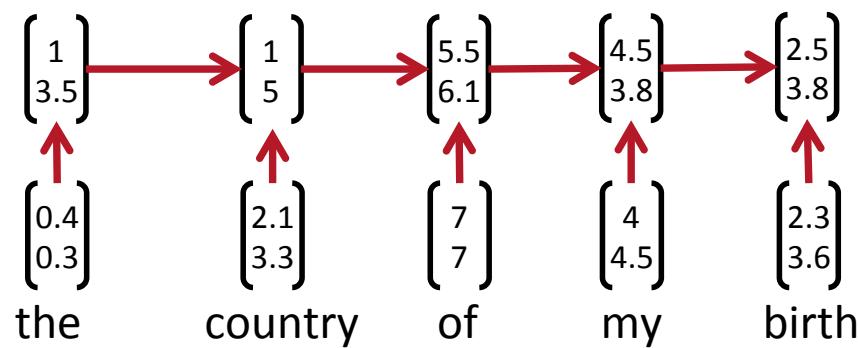
Learn jointly Structure and Representation

- Do we really need to learn this structure?
- Are languages structurally recursive?

Recursive vs Recurrent



- Topology of the recurrent networks is always a **chain**, which is a special kind of tree.

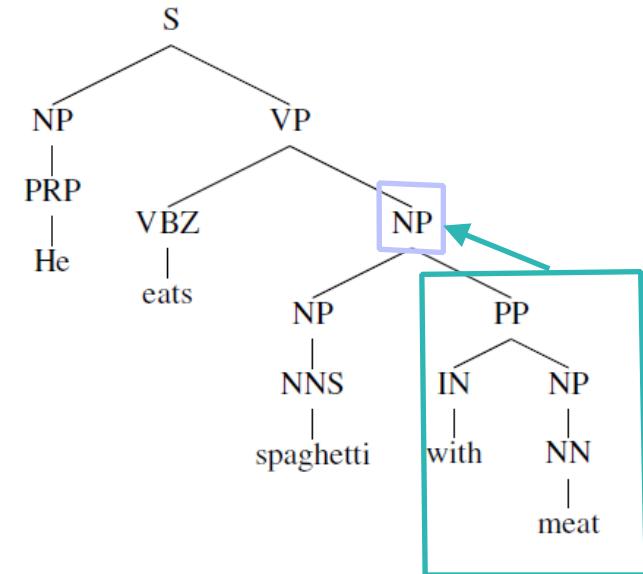
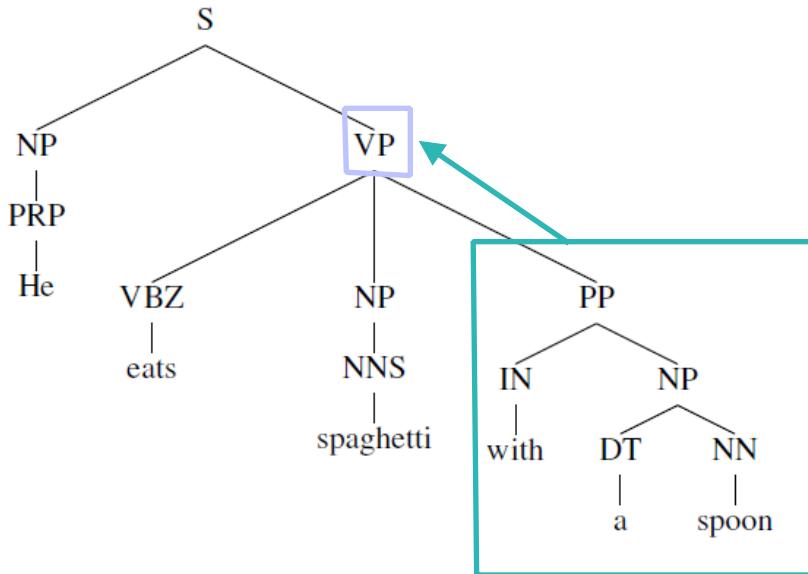


Are languages recursive?

- Debatable.
- Yet, recursion is helpful in describing natural language
- Example: “the church which has nice windows”, a noun phrase containing a relative clause that contains a noun phrase. (Kind of recursive statement)

Are languages recursive?

- Argument 1: Helpful in disambiguation



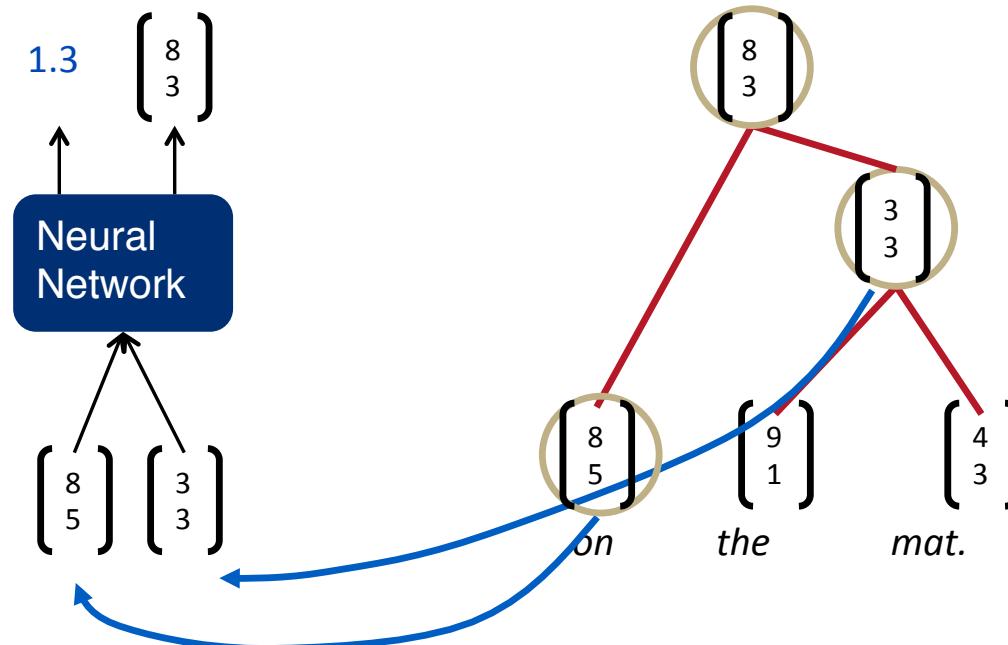
Are languages recursive?

- Argument 2: Helpful for some tasks (e.g **coreference resolution**) to refer to specific phrases:
 - John and Jane went to a big festival. **They** enjoyed **the trip** and the music **there**.
 - “they”: John and Jane
 - “the trip”: went to a big festival
 - “there”: big festival
- Argument 3: Labeling becomes less clear when specific to sub phrases
 - I liked the **bright screen** but not the **buggy slow keyboard** of the phone. **It** was a pain to type with. **It** was nice to look at.
- Argument 4: Works better for some tasks to use grammatical tree structure

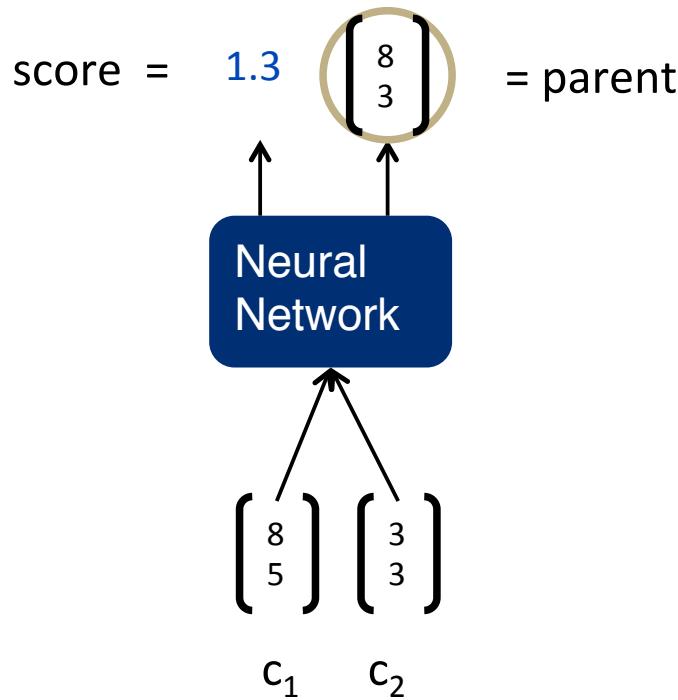
Still up for Debate!

Recursive Neural Networks for Structure Prediction

- Inputs: Two candidate children's representation
- Outputs:
 1. The semantic representation of the parent composition node
 2. Score of how plausible the new node could be

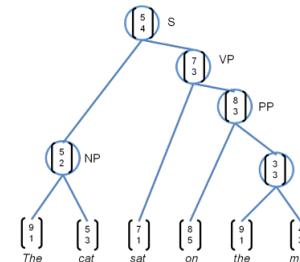


Definition of Recursive Neural Network

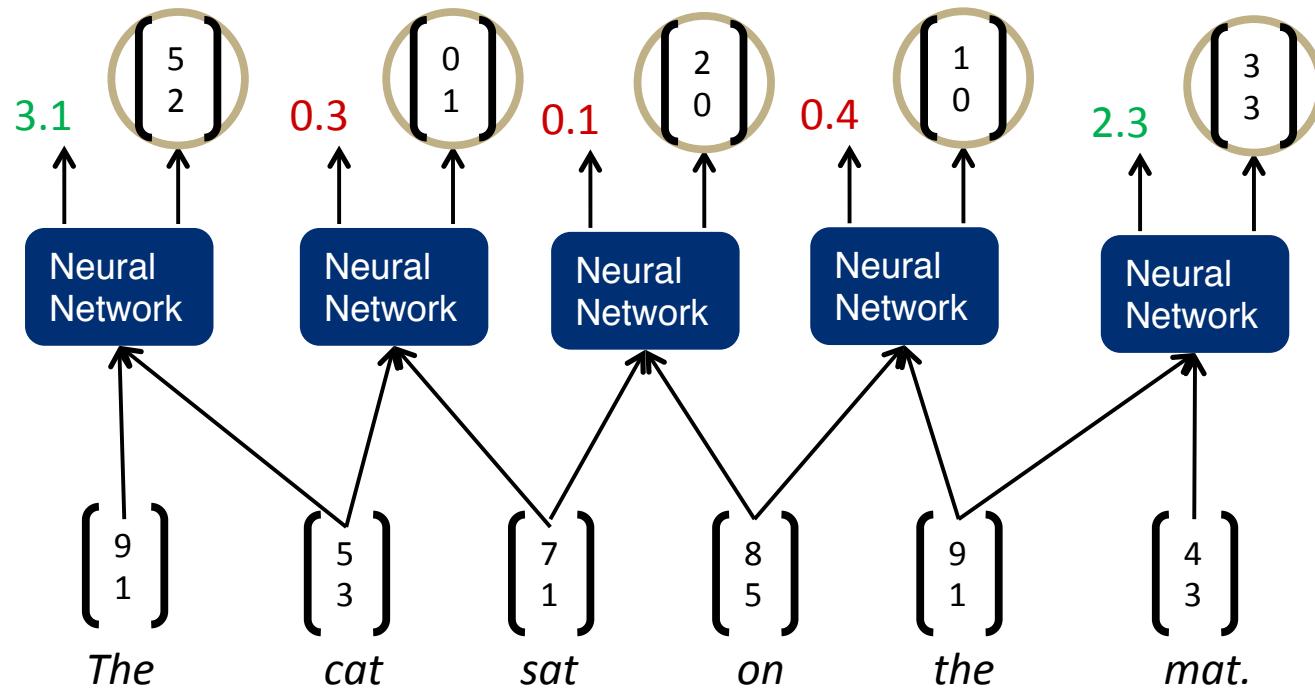


$$\left. \begin{aligned} \text{score} &= U^T p \\ p &= \tanh\left(W \begin{pmatrix} c_1 \\ c_2 \end{pmatrix} + b\right), \end{aligned} \right\}$$

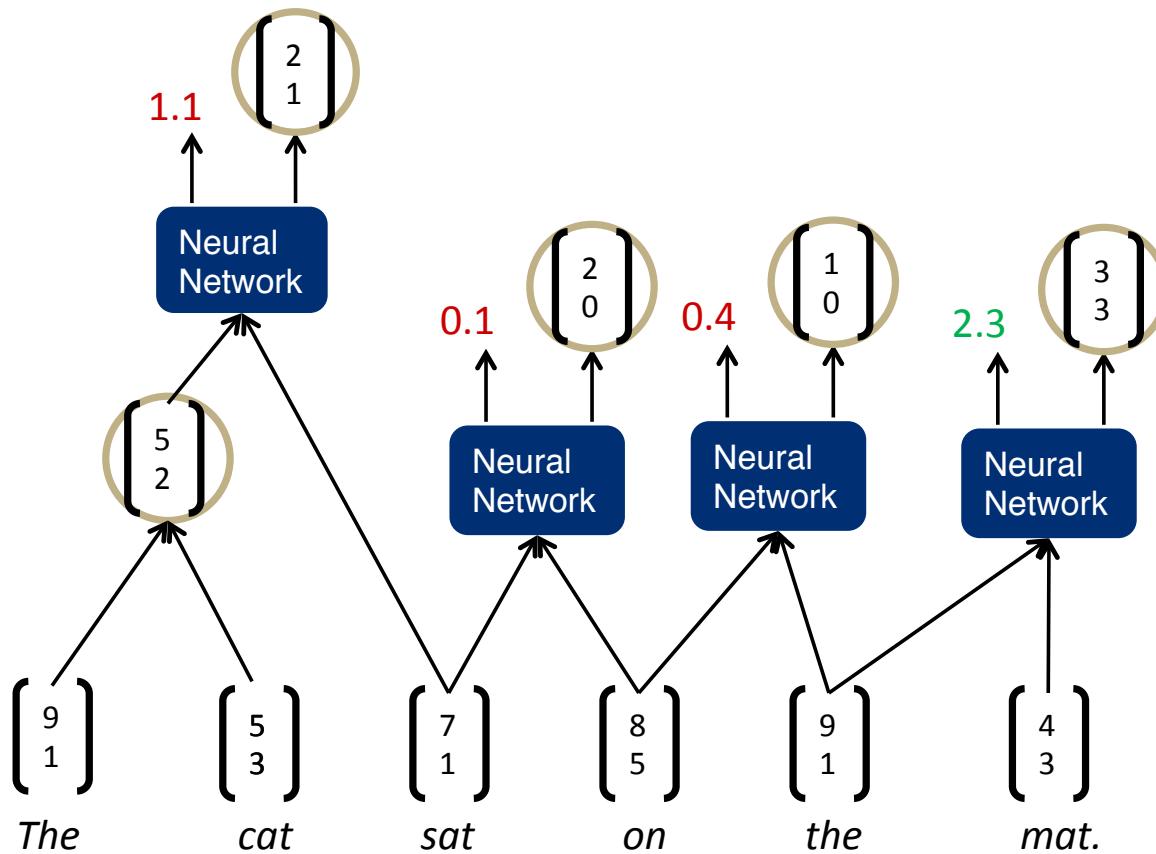
Same W parameters at all nodes of the tree



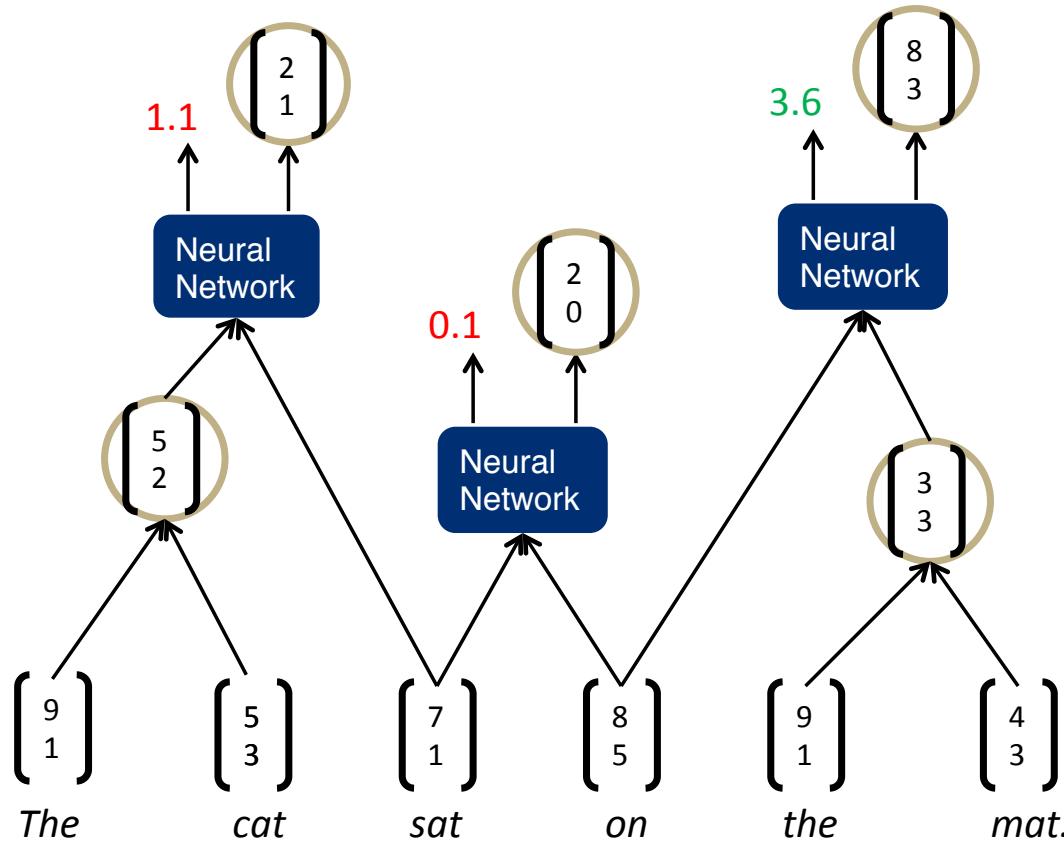
Parsing sentence with RNN



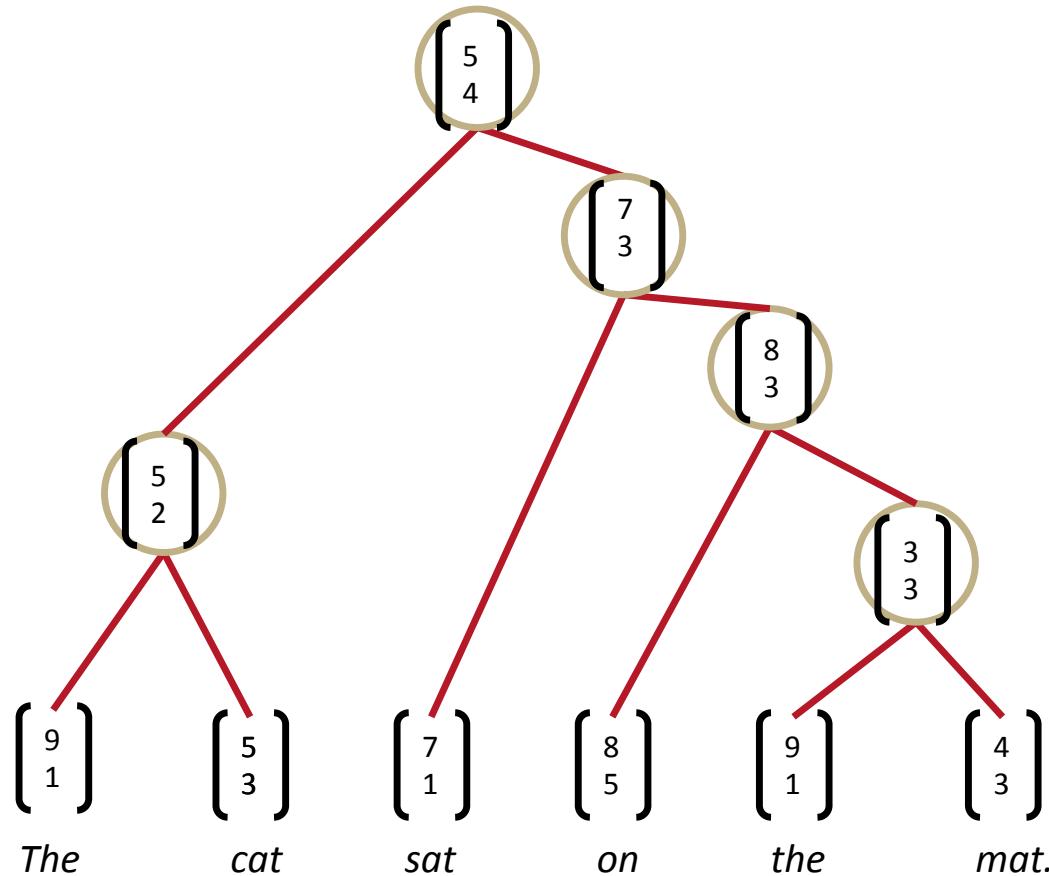
Parsing sentence with RNN



Parsing sentence with RNN



Parsing sentence with RNN



Score

- The score of a tree is computed by the sum of the parsing decision scores at each node:

$$s(x, y) = \sum_{n \in nodes(y)} s_n$$



Max-Margin Objective

- Similar to max-margin parsing (Taskar et al. 2004), a supervised max-margin objective

$$J = \sum_i s(x_i, y_i) - \max_{y \in A(x_i)} (s(x_i, y) + \Delta(y, y_i))$$

- The loss $\Delta(y, y_i)$ penalizes all incorrect decisions
- Structure search for $A(x)$ was maximally greedy
 - Instead: Beam Search with Chart

Backpropagation Through Structure

Introduced by Goller & Küchler (1996)

Principally the same as general backpropagation

$$\delta^{(l)} = \left((W^{(l)})^T \delta^{(l+1)} \right) \circ f'(z^{(l)}),$$

$$\frac{\partial}{\partial W^{(l)}} E_R = \delta^{(l+1)} (a^{(l)})^T + \lambda W^{(l)}$$

Three differences resulting from the recursion and tree structure:

1. Sum derivatives of W from all nodes
2. Split derivatives at each node
3. Add error messages

Backpropagation Through Structure

- Sum derivatives of all nodes

You can actually assume it's a different W at each node

Intuition via example:

$$\begin{aligned} & \frac{\partial}{\partial W} f(W(f(Wx))) \\ = & f'(W(f(Wx))) \left(\left(\frac{\partial}{\partial W} W \right) f(Wx) + W \frac{\partial}{\partial W} f(Wx) \right) \\ = & f'(W(f(Wx))) (f(Wx) + W f'(Wx)x) \end{aligned}$$

If we take separate derivatives of each occurrence, we get same:

$$\begin{aligned} & \frac{\partial}{\partial W_2} f(W_2(f(W_1x))) + \frac{\partial}{\partial W_1} f(W_2(f(W_1x))) \\ = & f'(W_2(f(W_1x))) (f(W_1x)) + f'(W_2(f(W_1x))) (W_2 f'(W_1x)x) \\ = & f'(W_2(f(W_1x))) (f(W_1x) + W_2 f'(W_1x)x) \\ = & f'(W(f(Wx))) (f(Wx) + W f'(Wx)x) \end{aligned}$$

Backpropagation Through Structure

- Split derivative at each node

During forward prop, the parent is computed using 2 children

$$\begin{array}{c}
 \begin{pmatrix} 8 \\ 3 \end{pmatrix} \\
 \swarrow \quad \searrow \\
 \begin{pmatrix} 8 \\ 5 \end{pmatrix} c_1 \quad \begin{pmatrix} 3 \\ 3 \end{pmatrix} c_2
 \end{array} \quad p = \tanh\left(\mathbf{w} \begin{pmatrix} c_1 \\ c_2 \end{pmatrix} + \mathbf{b}\right)$$

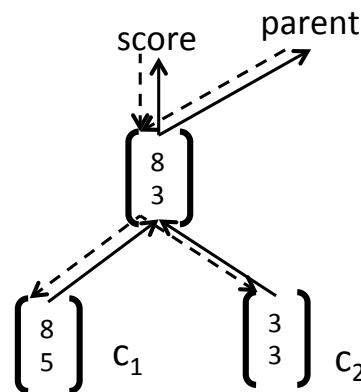
Hence, the errors need to be computed wrt each of them:

$$\begin{array}{c}
 \begin{pmatrix} 8 \\ 3 \end{pmatrix} \\
 \swarrow \quad \searrow \\
 \begin{pmatrix} 8 \\ 5 \end{pmatrix} c_1 \quad \begin{pmatrix} 3 \\ 3 \end{pmatrix} c_2
 \end{array} \quad \text{where each child's error is } n\text{-dimensional}$$

$$\delta_{p \rightarrow c_1 c_2} = [\delta_{p \rightarrow c_1} \delta_{p \rightarrow c_2}]$$

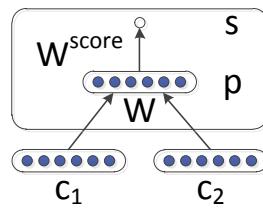
Backpropagation Through Structure

- Add error messages:
- At each node:
 - What came up (fprop) must come down (bprop)
 - Total error messages δ = error messages from parent + error message from own score



Too Simple?

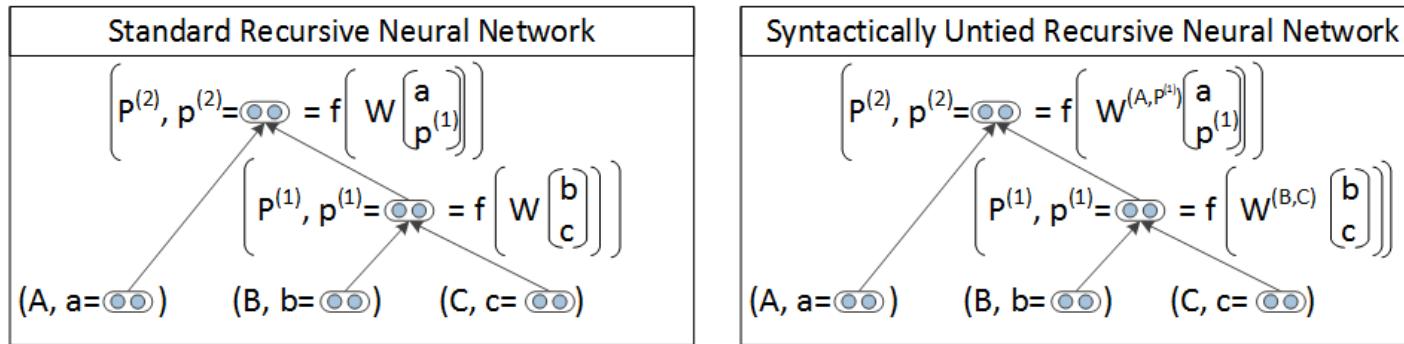
- Good results with single matrix RNN (more later)
- Single weight matrix RNN could capture some phenomena but not adequate for more complex, higher order composition and parsing long sentences



- The composition function is the same for all syntactic categories, punctuation, etc

SU-RNN: Syntactically Untied RNN

- Idea: Condition the composition function on the syntactic categories, “untie the weights”
- Allows for different composition functions for pairs of syntactic categories, e.g. Adv + AdjP, VP + NP
- Combines discrete syntactic categories with continuous semantic information



Experiments

Parser	Test, All Sentences
Stanford PCFG, (Klein and Manning, 2003a)	85.5
Stanford Factored (Klein and Manning, 2003b)	86.6
Factored PCFGs (Hall and Klein, 2012)	89.4
Collins (Collins, 1997)	87.7
SSN (Henderson, 2004)	89.4
Berkeley Parser (Petrov and Klein, 2007)	90.1
CVG (RNN) (Socher et al., ACL 2013)	85.0
CVG (SU-RNN) (Socher et al., ACL 2013)	90.4
Charniak - Self Trained (McClosky et al. 2006)	91.0
Charniak - Self Trained-ReRanked (McClosky et al. 2006)	92.1

- 3.8% higher F1
- 20% faster than Stanford Factored parser

Analysis of resulting phrase vectors

All the figures are adjusted for seasonal variations

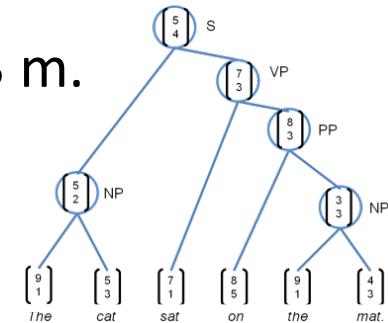
1. All the numbers are adjusted for seasonal fluctuations
2. All the figures are adjusted to remove usual seasonal patterns

Knight-Ridder wouldn't comment on the offer

1. Harsco declined to say what country placed the order
2. Coastal wouldn't disclose the terms

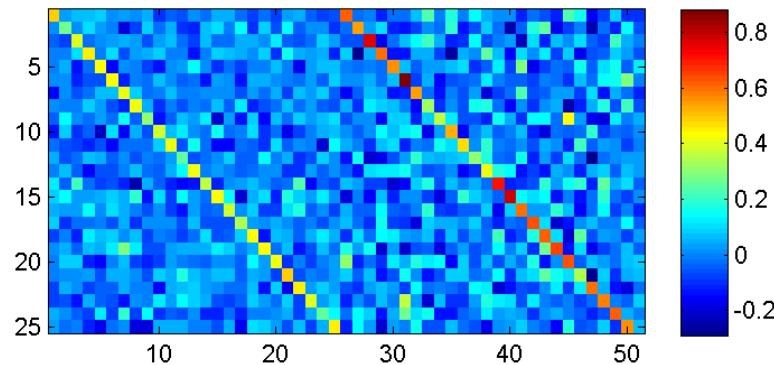
Sales grew almost 7% to \$UNK m. from \$UNK m.

1. Sales rose more than 7% to \$94.9 m. from \$88.3 m.
2. Sales surged 40% to UNK b. yen from UNK b.

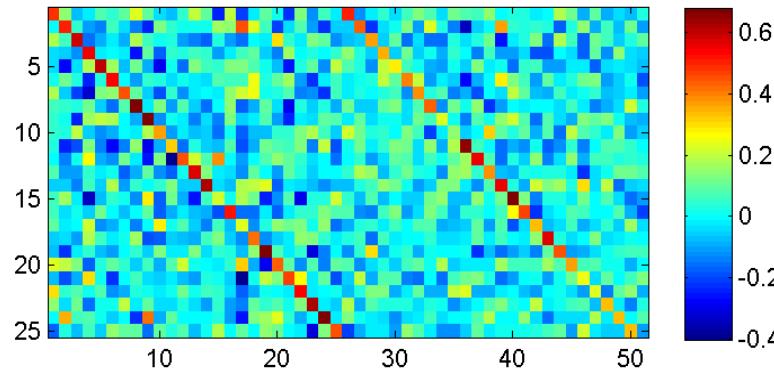


SU-RNN Analysis

- Learns notion of soft head words



DT-NP
the cat



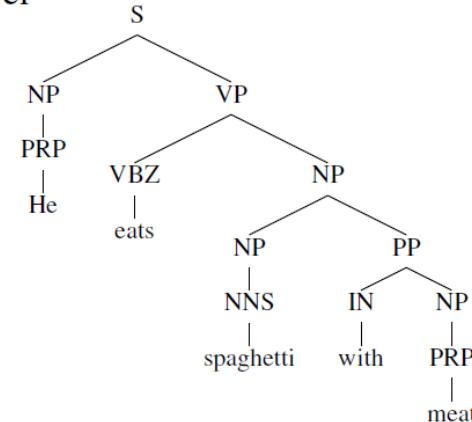
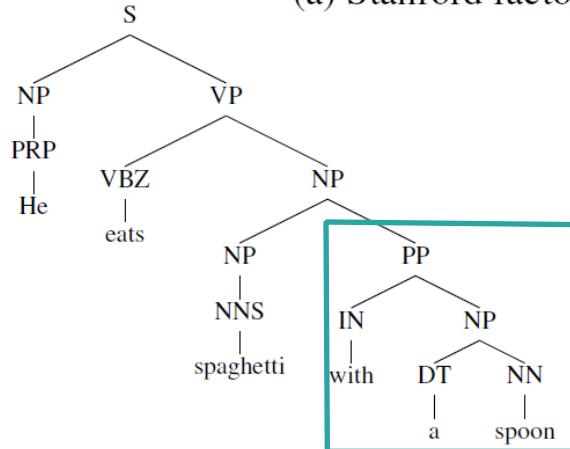
VP-NP
read book

SU-RNN Analysis

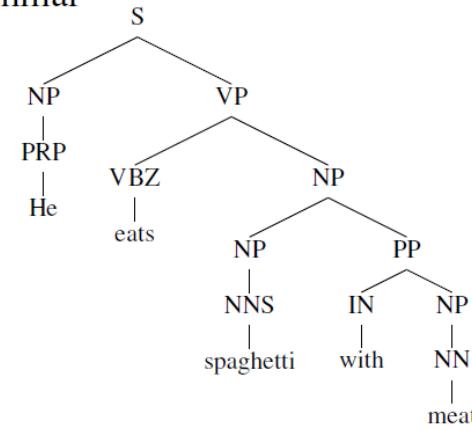
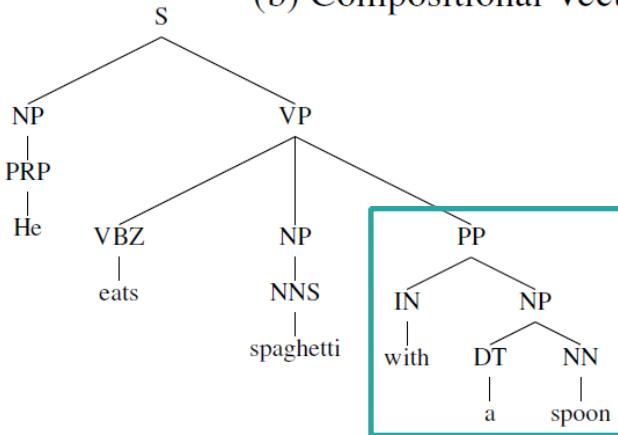
- Can transfer semantic information from single related example
- Train sentences:
 - He eats spaghetti with a fork.
 - She eats spaghetti with pork.
- Test sentences
 - He eats spaghetti with a spoon.
 - He eats spaghetti with meat.

SU-RNN Analysis

(a) Stanford factored parser



(b) Compositional Vector Grammar

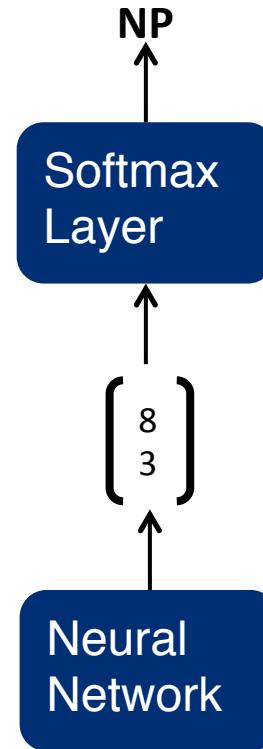


Labeling tasks in RNN

- We can just use softmax classifier over the representation vector of each node:

$$p(c|p) = \text{softmax}(Sp)$$

- Training similarly with standard cross-entropy error + scores



Outline

- Motivation
- Recursive Neural Networks
 - Compositionality
 - Structure Prediction: Parsing
 - Backpropagation Through Structure
 - Minor twist leads to an improved Parser
- Matrix-Vector RNNs
- Recursive Neural Tensor Networks

Compositionality By Matrix-Vector RNNs

$$p = \tanh\left(W \begin{bmatrix} c_1 \\ c_2 \end{bmatrix} + b\right)$$

One way to make the composition function more powerful was by untying the weights W

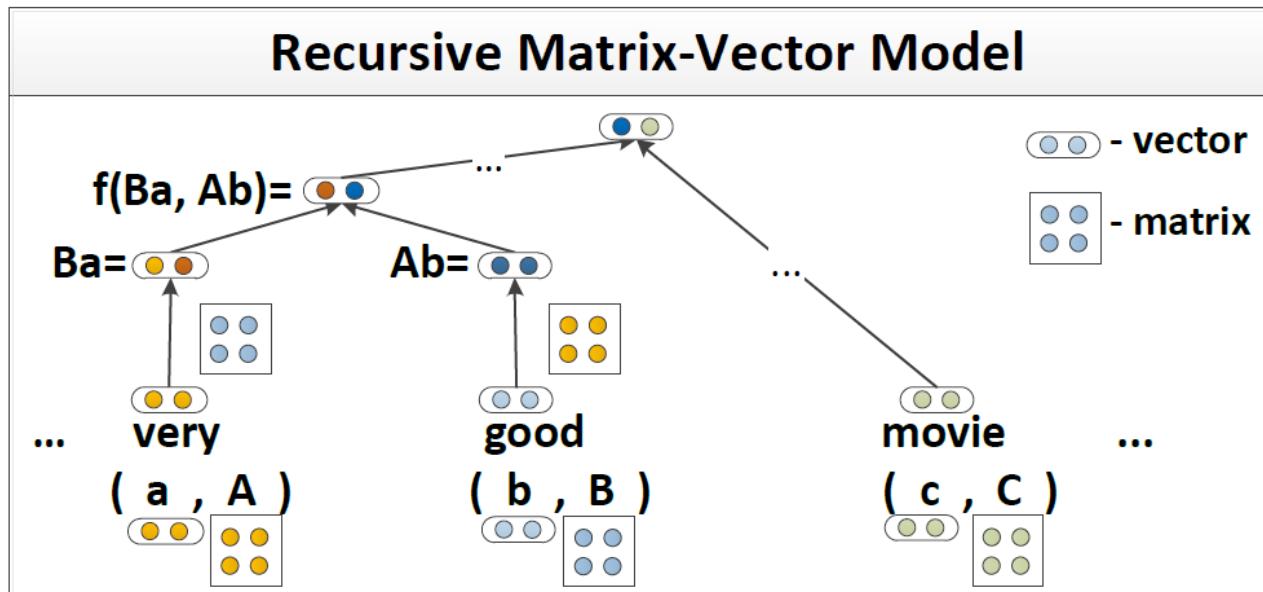
But what if words act mostly as an operator, e.g. “very” in
very good

Proposal: A new composition function

Compositionality By Matrix-Vector RNNs

$$p = \tanh\left(W \begin{bmatrix} c_1 \\ c_2 \end{bmatrix} + b\right)$$

$$p = \tanh\left(W \begin{bmatrix} c_2 c_1 \\ c_1 c_2 \end{bmatrix} + b\right)$$



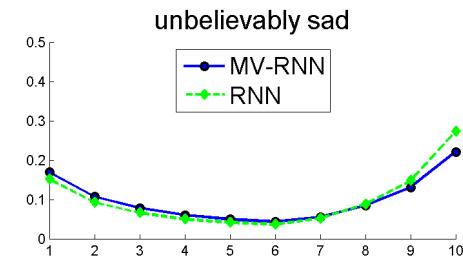
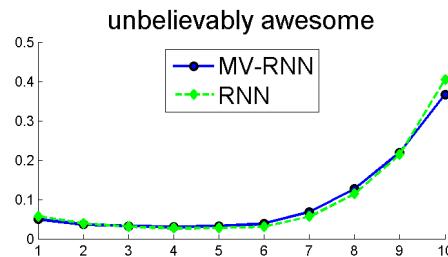
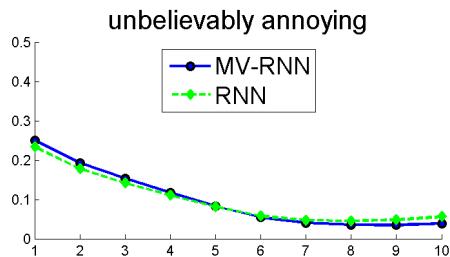
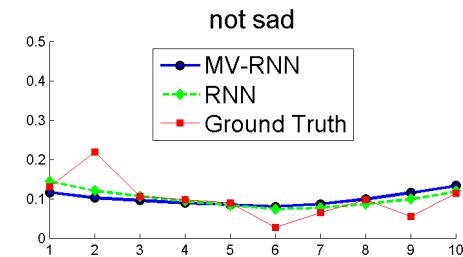
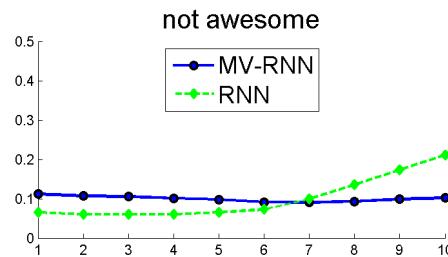
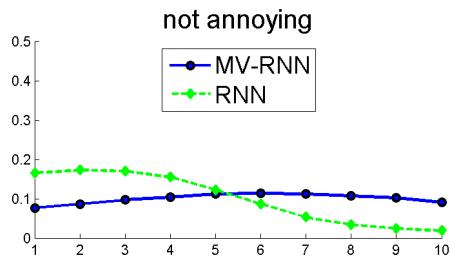
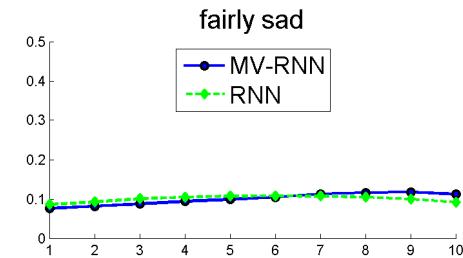
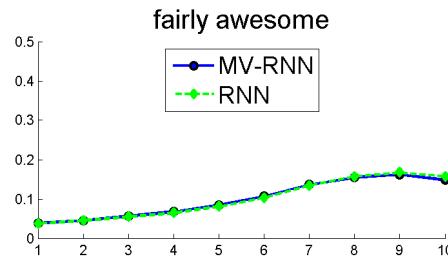
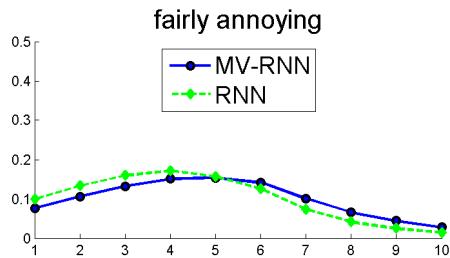
Predicting Sentiment Distributions

Method	Avg KL
Uniform	0.327
Mean train	0.193
$p = \frac{1}{2}(a + b)$	0.103
$p = a \otimes b$	0.103
$p = [a; b]$	0.101
$p = Ab$	0.103
RNN	0.093
Linear MVR	0.092
MV-RNN	0.091

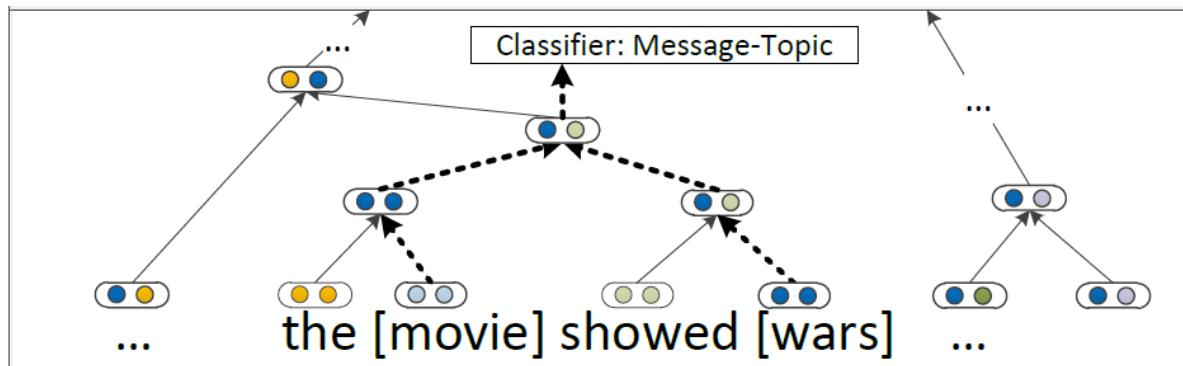
Method	Acc.
Tree-CRF (Nakagawa et al., 2010)	77.3
RAE (Socher et al., 2011c)	77.7
Linear MVR	77.1
MV-RNN	79.0

Predicting Sentiment Distributions via MV-RNN

Good example for non-linearity in language



Relationship Classification with MV-RNN



Relationship	Sentence with labeled nouns for which to predict relationships
Cause-Effect(e_2, e_1)	Avian [influenza] $_{e_1}$ is an infectious disease caused by type a strains of the influenza [virus] $_{e_2}$.
Entity-Origin(e_1, e_2)	The [mother] $_{e_1}$ left her native [land] $_{e_2}$ about the same time and they were married in that city.
Message-Topic(e_2, e_1)	Roadside [attractions] $_{e_1}$ are frequently advertised with [billboards] $_{e_2}$ to attract tourists.

Classifier	Feature Sets	F1
SVM	POS, stemming, syntactic patterns	60.1
SVM	word pair, words in between	72.5
SVM	POS, WordNet, stemming, syntactic patterns	74.8
SVM	POS, WordNet, morphological features, thesauri, Google n -grams	77.6
MaxEnt	POS, WordNet, morphological features, noun compound system, thesauri, Google n -grams	77.6
SVM	POS, WordNet, prefixes and other morphological features, POS, dependency parse features, Levin classes, PropBank, FrameNet, NomLex-Plus, Google n -grams, paraphrases, TextRunner	82.2
RNN	-	74.8
Lin.MVR	-	73.0
MV-RNN	-	79.1
RNN	POS,WordNet,NER	77.6
Lin.MVR	POS,WordNet,NER	78.7
MV-RNN	POS,WordNet,NER	82.4

Outline

- Motivation
- Recursive Neural Networks
 - Compositionality
 - Structure Prediction: Parsing
 - Backpropagation Through Structure
 - Minor twist leads to an improved Parser
- Matrix-Vector RNNs
- **Recursive Neural Tensor Networks**

Motivation - Sentiment Detection

Most methods start with a bag of words
+ linguistic features/processing/lexica

But such methods (including tf-idf) can't
distinguish:

Difficult!

- + white blood cells destroying an infection
- an infection destroying white blood cells

Motivation - Sentiment Detection

- For dataset of single sentence movie reviews (Pang and Lee, 2005) accuracy never reached above 80% for >7 years
- Harder cases require actual understanding of **negation and its scope** + other semantic effects

Motivation - Sentiment Detection

Examples from Movie Reviews Dataset:

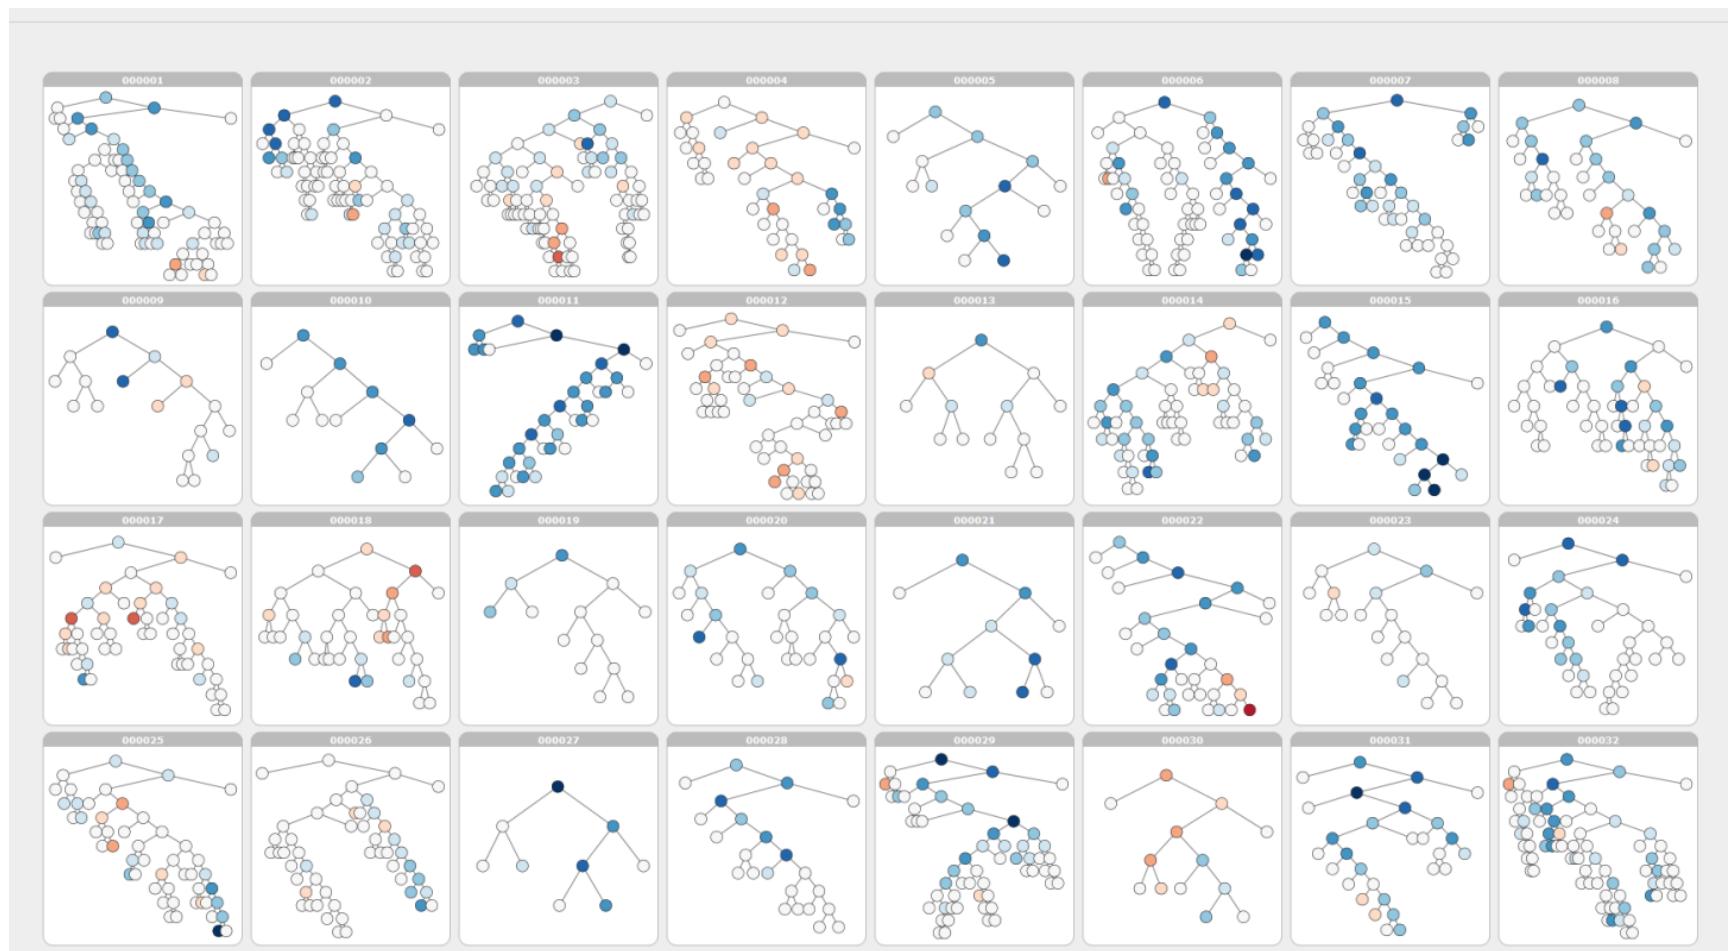
Stealing Harvard doesn't care about cleverness, wit or any other kind of intelligent humor.

There are slow and repetitive parts but it has just enough spice to keep it interesting.

What is missing?

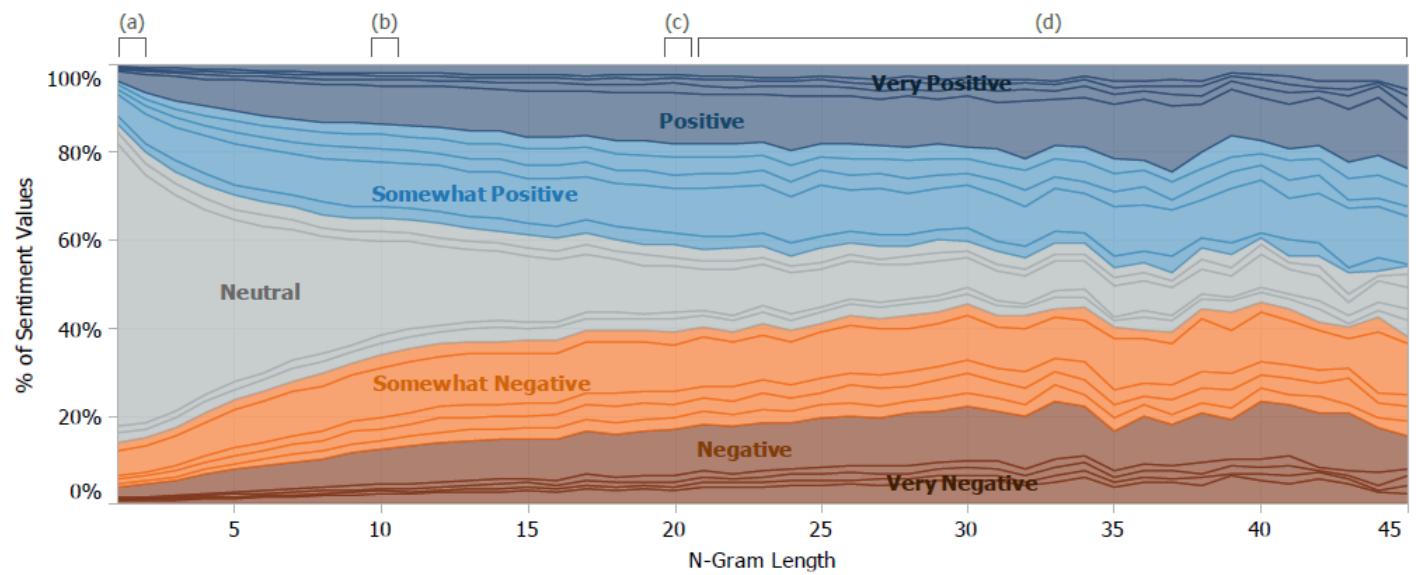
- 1. Compositional Training Data**
- 2. Better Compositional model**

Solution 1: New Sentiment Treebank



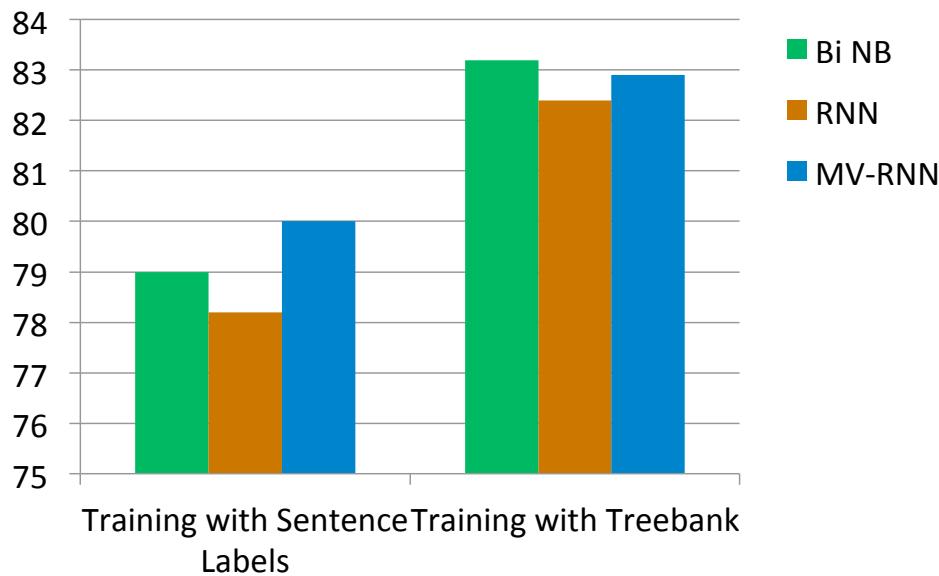
Solution 1: New Sentiment Treebank

- Parse trees of 11,855 sentences
- 215,154 phrases with labels
- Allows training and evaluating with compositional information



Using Better Dataset Helped All Models

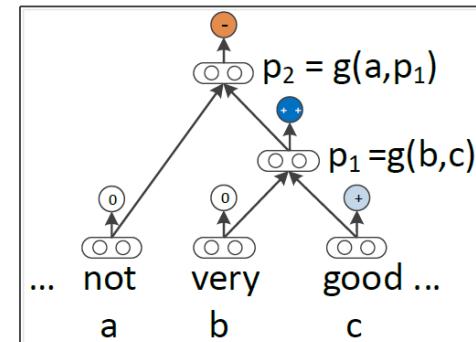
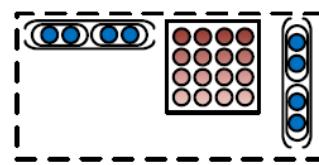
- Binary(pos/neg) full sentence classification



- Yet, hard negation cases are still mostly incorrect
- Better/stronger compositional model

Solution 2: New Compositional Model

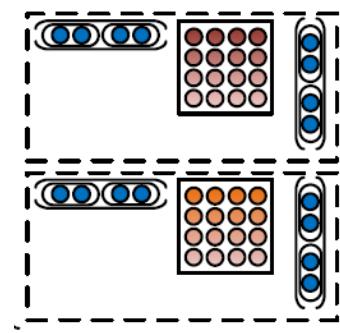
- Recursive Neural Tensor Network
- More expressive than previous RNNs
- Idea: Allow more interactions of vectors



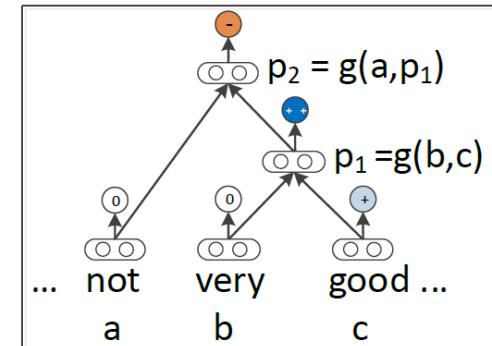
$$\begin{bmatrix} b \\ c \end{bmatrix}^T V \quad \begin{bmatrix} b \\ c \end{bmatrix}$$

Solution 2: New Compositional Model

- Recursive Neural Tensor Network

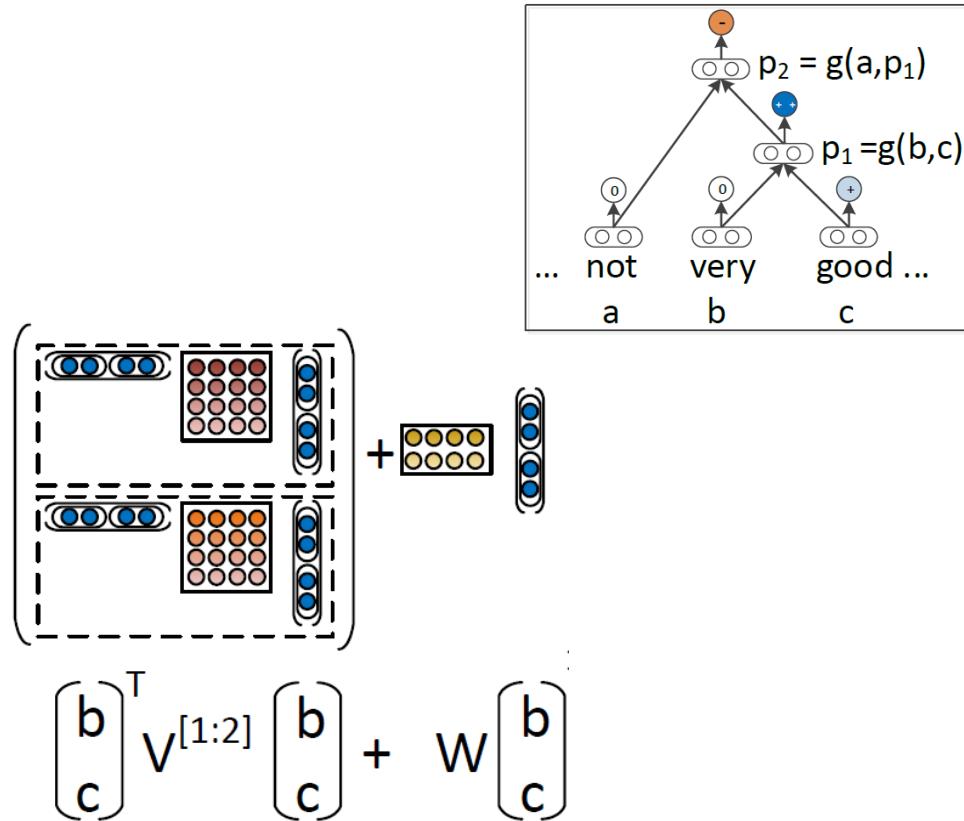


$$\begin{bmatrix} b \\ c \end{bmatrix}^T V^{[1:2]} \begin{bmatrix} b \\ c \end{bmatrix}$$



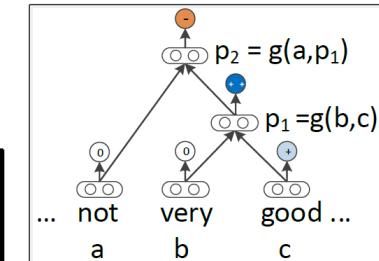
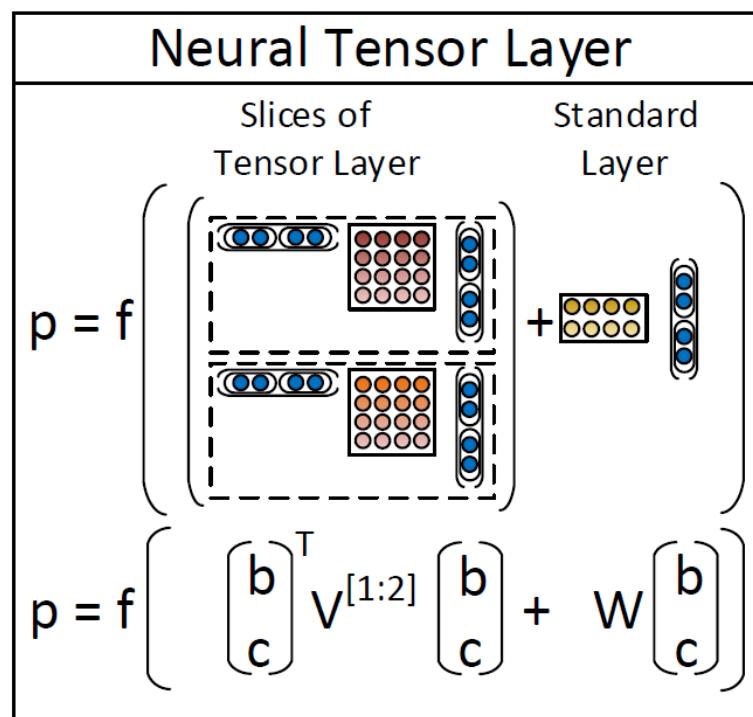
Solution 2: New Compositional Model

- Recursive Neural Tensor Network



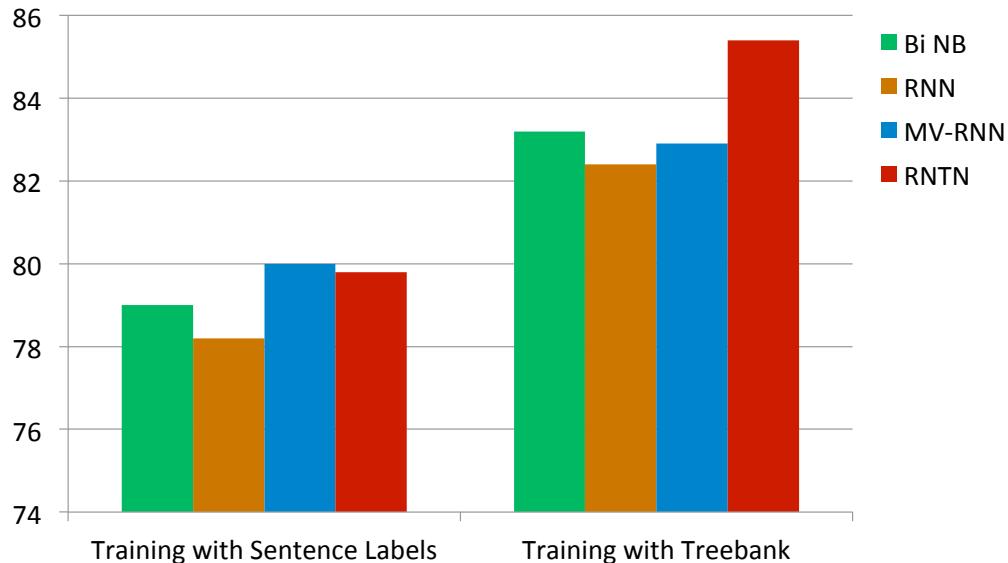
Recursive Neural Tensor Network

Recursive Deep Models for Semantic Compositionality Over a Sentiment Treebank
Socher et al. 2013



Positive/Negative Results on Treebank

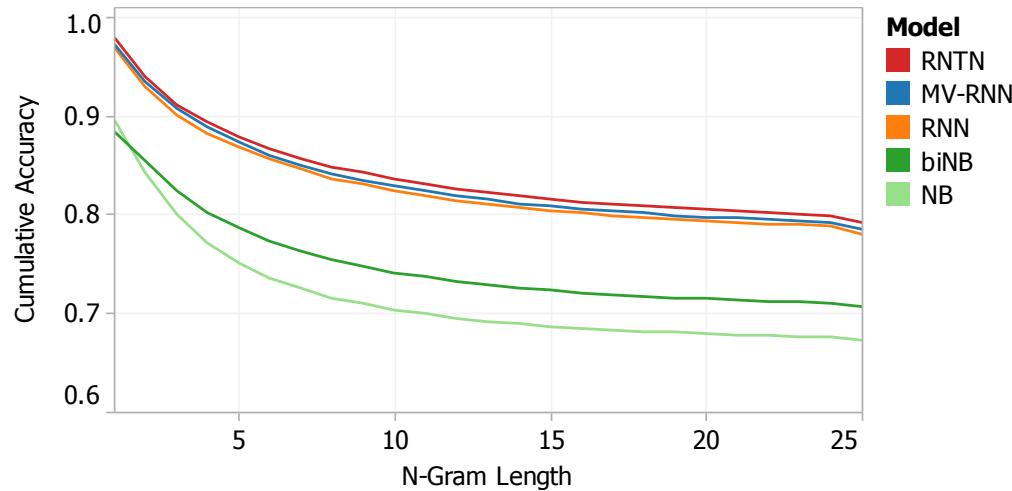
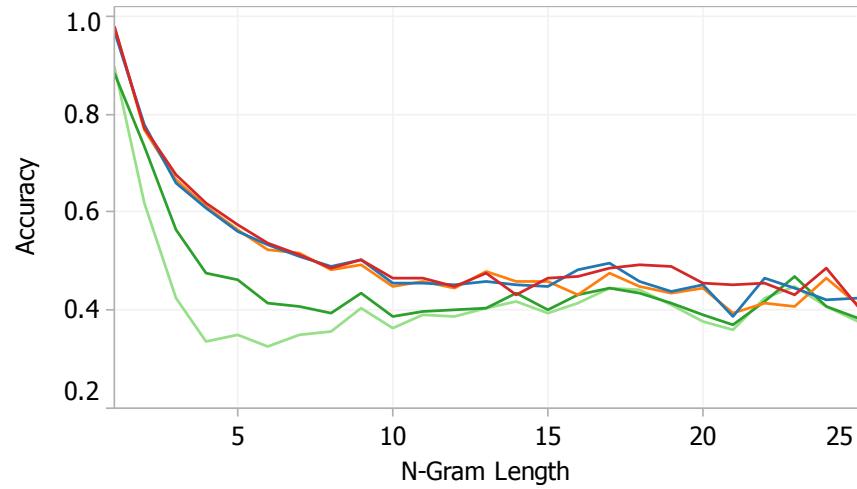
Classifying Sentences: Accuracy improves to 85.4



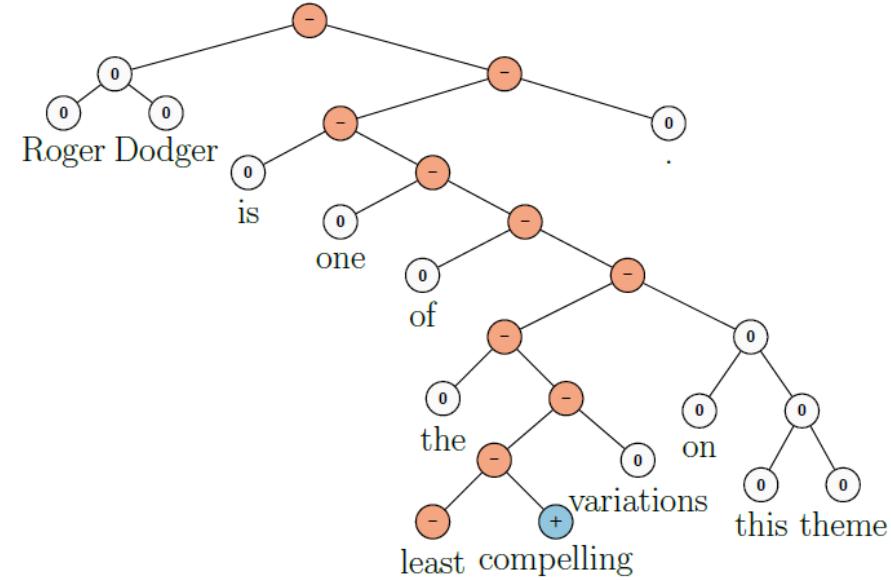
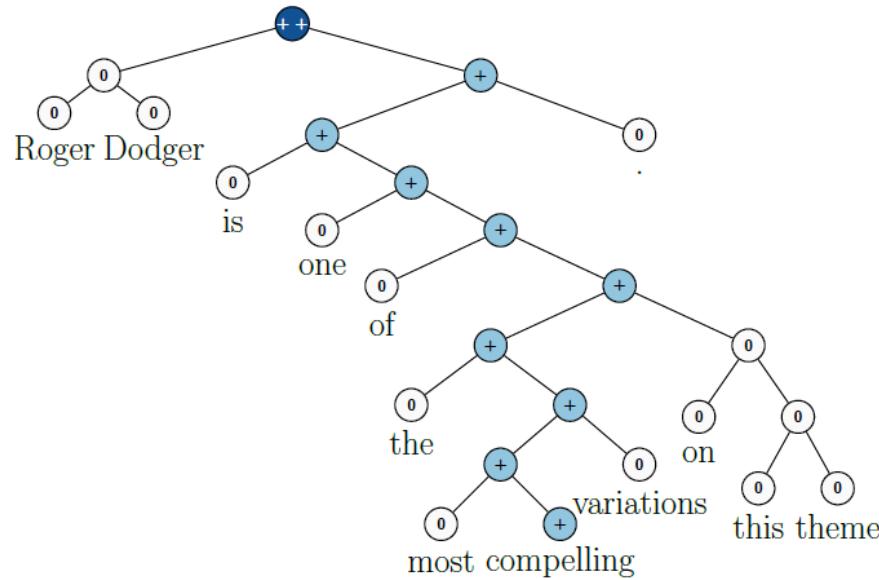
Model	Fine-grained		Positive/Negative	
	All	Root	All	Root
NB	67.2	41.0	82.6	81.8
SVM	64.3	40.7	84.6	79.4
BiNB	71.0	41.9	82.7	83.1
VecAvg	73.3	32.7	85.1	80.1
RNN	79.0	43.2	86.1	82.4
MV-RNN	78.7	44.4	86.8	82.9
RNTN	80.7	45.7	87.6	85.4

Root: Sentence Level

Final Grained Sentiment Classification Results



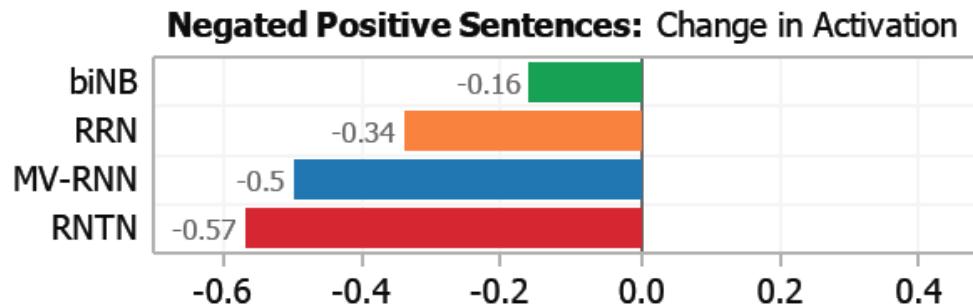
Negation Results



Negation Results

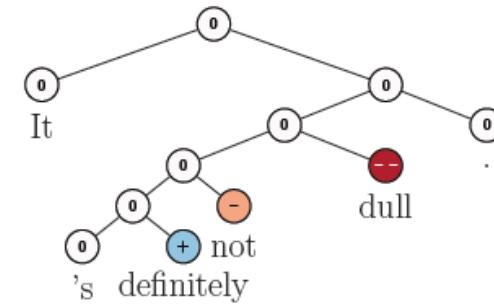
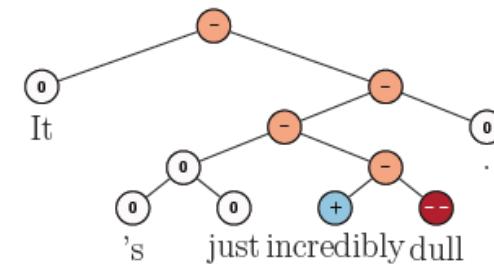
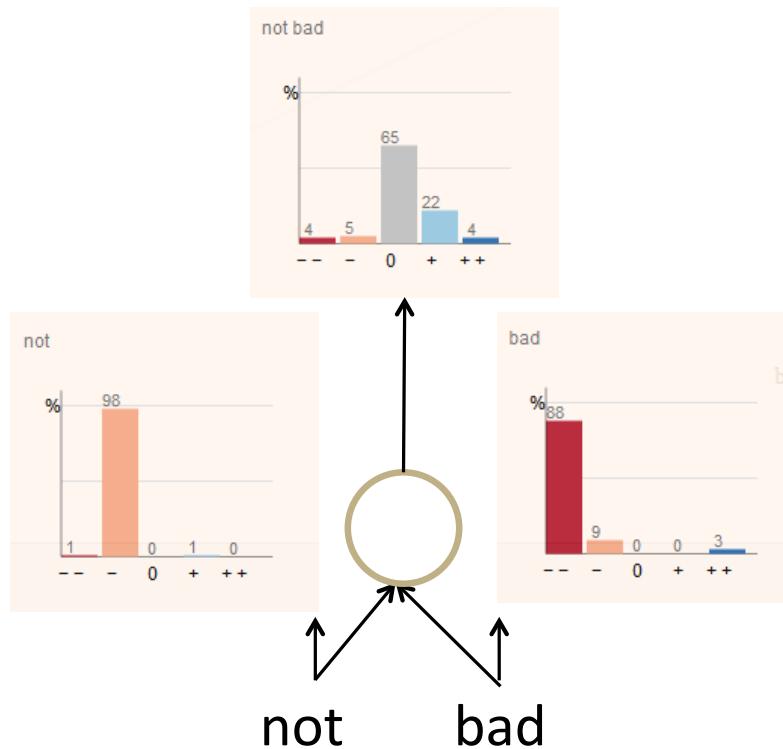
- Most methods capture that negation often makes things more negative (See Potts, 2010)
- Analysis on negation dataset
- Accuracy:

	Negated Positive
biNB	19.0
RNN	33.3
MV-RNN	52.4
RNTN	71.4



Results on Negating Negatives

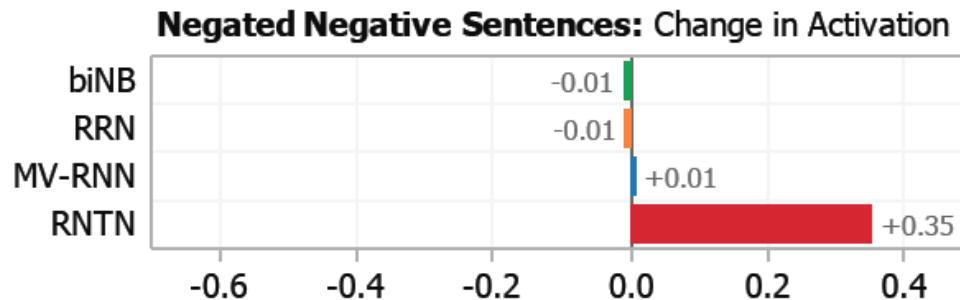
- But how about negating negatives?
- No flips, but positive activation should increase!



Results on Negating Negatives

- Evaluation: Positive activation should increase

Model	Accuracy	
	Negated Positive	Negated Negative
biNB	19.0	27.3
RNN	33.3	45.5
MV-RNN	52.4	54.6
RNTN	71.4	81.8



References

1. Learning Continuous Phrase Representations and Syntactic Parsing with Recursive Neural Networks (Socher et al. NIPS 2010)
2. Parsing Natural Scenes and Natural Language with Recursive Neural Networks (Socher et al. ICML 2011)
3. Semantic Compositionality through Recursive Matrix-Vector Spaces (Socher et al. EMNLP 2012)
4. Parsing with Compositional Vector Grammars (Socher et al. ACL 2013)
5. Recursive Deep Models for Semantic Compositionality Over a Sentiment Treebank (Socher et al. EMNLP 2013)

